



# Accelerating Data-Centers using NVMe and CUDA

Stephen Bates, PhD  
Technical Director, CSTO,  
PMC-Sierra



## Project Donard @ PMC-Sierra

- Donard is a PMC-Sierra CTO project that leverages NVM Express (NVMe) to accelerate data-center applications.
- It uses NVMe and (Remote) Direct Memory Access to enable the “Trifecta” of compute, network and storage PCIe devices.
- Builds on work done by Nvidia, Mellanox and others.

# The Problem: x86 Status Quo

Low parallelism = CPU / Memory / fabric saturation

<10K IOPS per CPU<sup>1</sup>

~1M IOPS per flash device

Fast Data  
workloads  
requiring deep  
parallelism  
are not efficient with  
x86



Fabric



~2M IOPS per 40G Link

**Fabric becomes overloaded**

<sup>1</sup> Assuming computation per IOP is high (e.g. image search, encryption, audio processing).

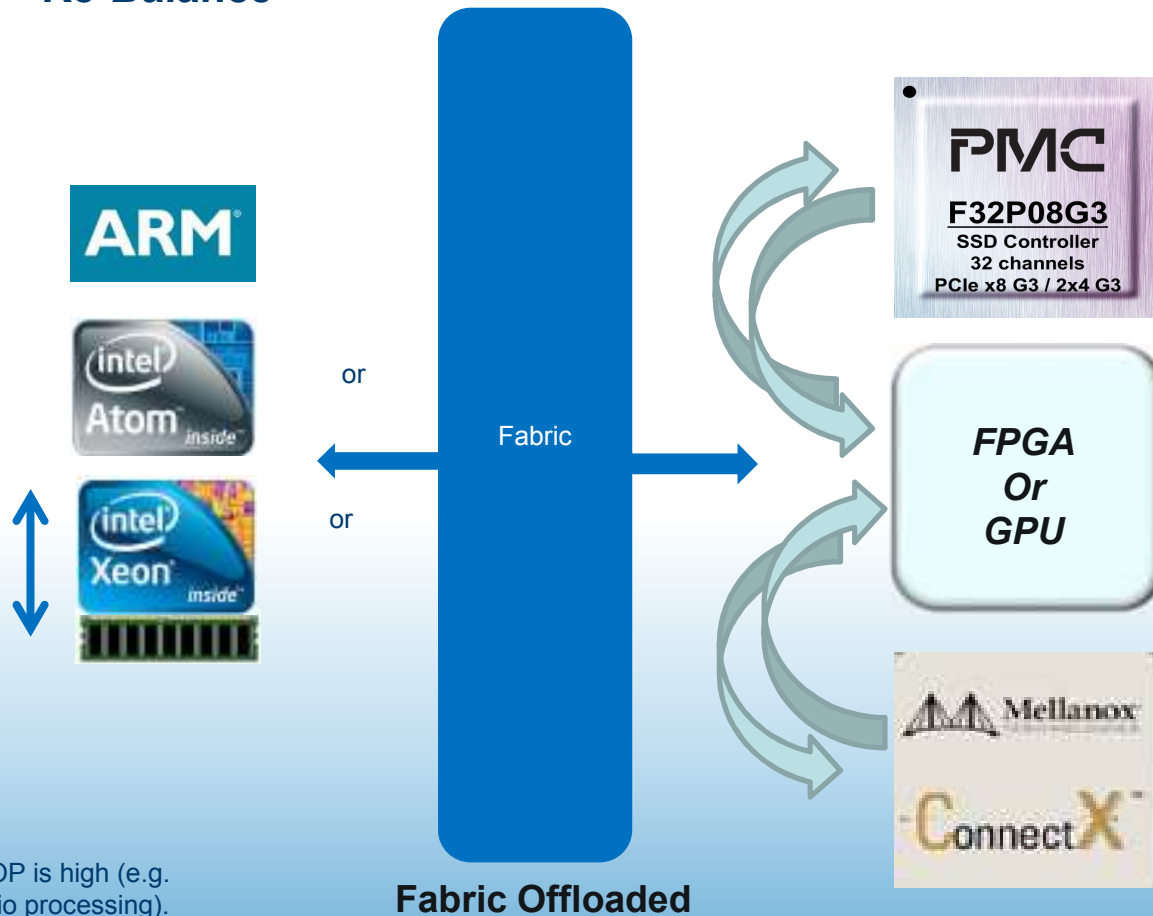
# The Solution: Donard

Pre-process algorithms in the data path

Re-Balance

A >99% filter rate drops IOPs to 10K.

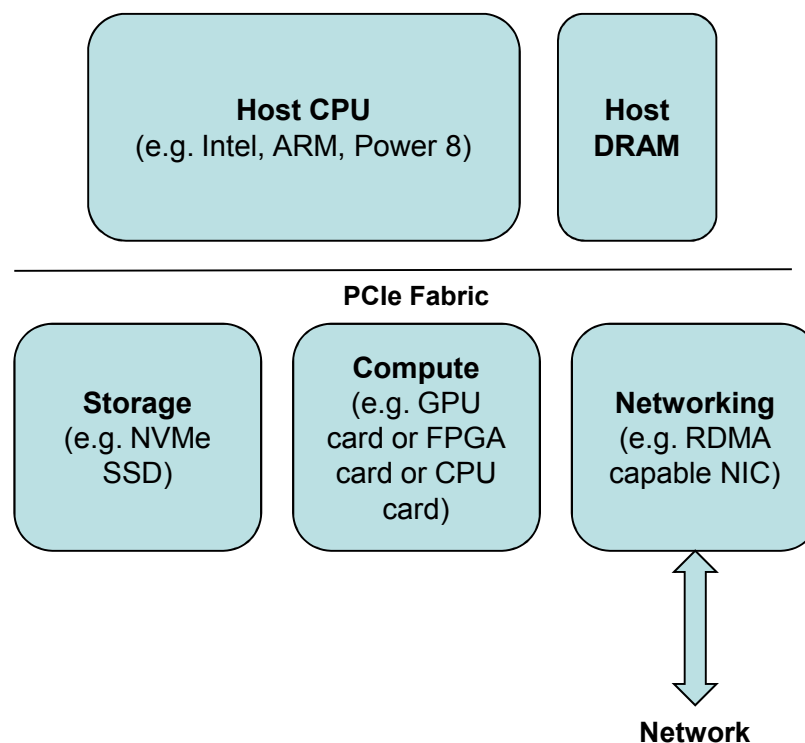
Simple Control  
And Management



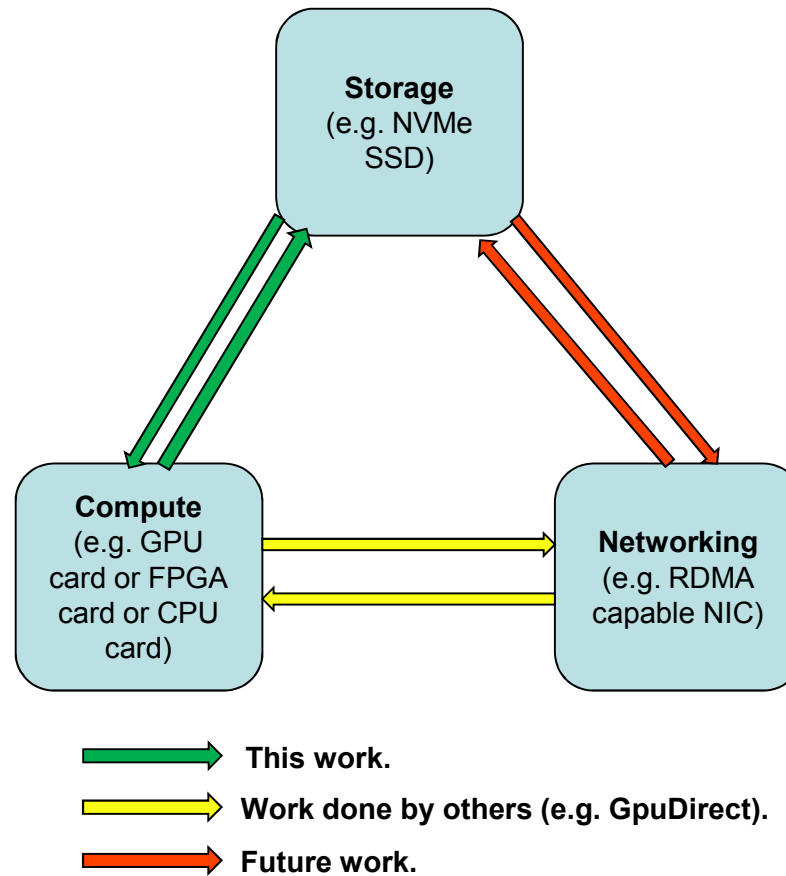
<sup>1</sup> Assuming computation per IOP is high (e.g. image search, encryption, audio processing).

## The PCIe “Trifecta”

- We want to enable the Trifecta on a PCIe fabric.
- This work leverages NVMe to enable two edges of this Trifecta.
- Others have worked on other arms (see next slide).



# The PCIe “Trifecta”







## Why NVMe Express?

- NVMe Express (NVMe) provides a consistent, open-source interface into PCIe storage devices.
- An NVMe driver has been part of both Linux and Windows Server for quite some time now.
- NVMe is extendible and is the focus of optimization within kernel.org<sup>1</sup>.
- Using NVMe makes Donard much more scalable and amenable to community development.
- The work in this paper should be applicable to any NVMe compliant drive.

<sup>1</sup>[http://kernelnewbies.org/Linux\\_3.13#head-3e5f0c2bcebc98efd197e3036dd814eadd62839c](http://kernelnewbies.org/Linux_3.13#head-3e5f0c2bcebc98efd197e3036dd814eadd62839c)

# Storage->Compute

- Built a Linux server running kernel 3.13.
- Installed an NVMe SSD and a Nvidia Tesla K20c.
- Modified the NVMe module in the kernel to add a new IOCTL that use DMA between SSD and the GPU card.
- Used CUDA 6.0 Peer-To-Peer (p2p) APIs to enable the DMA.
- Measured the impact of the new IOCTL on bandwidth and host DRAM utilization.



Technique	Bandwidth <sup>1</sup> (GB/s)	DRAM Utilization <sup>2</sup>
<b>Classical</b>	1.9	5230.0
<b>Donard (DMA)</b>	2.5	1.0

<sup>1</sup> Bandwidth was measured on our server which had a very standard PCIe fabric using a total transfer size of GB. Tests run 10 times. Results may vary depending on your PCIe architecture.

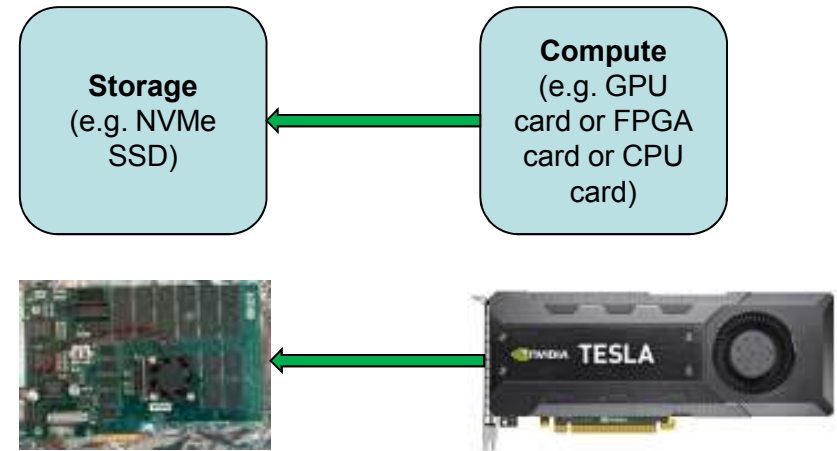
<sup>2</sup> DRAM utilization estimated using the page fault counters in the x86 CPU. Normalized to Donard performance.



# Compute->Storage

- Similar IOCTL as previous slide with direction changed.
- Had to resolve the file extents issue. Several options:
  1. Overwrite an existing file.
  2. Create a file using existing OS constructs (slow).
  3. Modify the file stats properties to allow uninitialized extents – possible security issue<sup>1</sup>.
- Measured the impact of the new IOCTL on bandwidth and host DRAM utilization.
- Still trying to determine why DMA method is slower. Suspect issue with PCIe architecture in server.

<sup>1</sup><https://lwn.net/Articles/492959/>



Technique	Bandwidth <sup>1</sup> (GB/s)	DRAM Utilization <sup>2</sup>
<b>Classical</b>	1.51	6012.0
<b>Donard (DMA)</b>	0.65	1.0

<sup>1</sup> Bandwidth was measured on our server which had a very standard PCIe fabric using a total transfer size of GB. Tests run 10 times. Results may vary depending on your PCIe architecture.

<sup>2</sup> DRAM utilization estimated using the page fault counters in the x86 CPU. Normalized to Donard performance.



## Donard in the Data-Center

- Data-Centers (DC) are deploying flash in a direct attach model to aid in application acceleration.
- Although SATA is prevalent today many DCs see a shift to PCIe attached flash and like the idea of using NVMe.
- Some DC customers are offloading applications to heterogeneous compute platforms such as GPU cards and FPGA cards<sup>1</sup>.
- Donard assists in this offload and also reduces the burden on the host processor and host DRAM.

# Donard DC Application - Haystack

- We wrote a program to search for the PMC logo in a large (10,000+) image database.
- Performance improved as we migrated to DMA on a SSD+GPU compared to a traditional solution.
- Note it also moves the bottleneck from the host DRAM interface to the GPU.
- Other applications might include sorting and write-caching.



	HDD	SDD	
	Mpix/s	Mpix/s	Bottleneck
CPU	77.0	122.8	CPU
CUDA	95.1	312.5	DRAM
<b>DONARD</b>	<b>N/A</b>	<b>534.2</b>	<b>GPU</b>



## Next Steps

- Determine reason for slow write performance from GPU into SSD.
- Add RDMA capable NICs to the Donard platform and complete the “Trifecta”, this is work in progress.
- Enable Donard within the community so others can benefit and contribute (GitHub?). Already working with Steve Swanson’s team at UCSD.
- Work with the NVM Express standards body to incorporate some of the Donard ideas into NVMe.



# Thank You

